ETH tech

F. Vogel

HPCSE I HS16

12-929-266

General

→ setxkbmap ch -variant de_nodeadkeys

Makefile make

```
CXX = g++
CFLAGS = -Wall -O3 -std=c++11

test: test.cpp
	$(CXX) $(CFLAGS) -o $@ $<

clean:
	rm -f test
```

Threads

compile g++ -std=c++11 -pthread

import #include <thread>

launch func std::thread t(foo, arg1);

launch λ std::thread([&]() { //do something });

mutex:

import #include <mutex>

declare std::mutex mtx;

lock
```
mtx.lock();
  ⋮
mtx.unlock();
```

lock_guard std::lock_guard<std::mutex> l(mtx);
↦ locked when initialized
↦ unlocked when destroyed

unique_lock std::unique_lock<std::mutex> l(mtx1, mtx2);
↦ locks work with multiple locks (no deadlock)
l.unlock(); l.lock(); for manual

Example

```
int nthreads = 4;
int nstep = N/nthreads;
std::vector<std::thread> threads(nthreads);
for (int t=0; t<nthreads; t++) {
  threads[t] = thread([&,t] {
    for (int i = t*nstep; i<(t+1)*nstep; i++){
      z[i] = x[i] + y[i];
    });
}
for (std::thread& t : threads) {
  t.join();
}
```

MPI

compile mpicc, mpic++

import #include <mpi.h>

run mpiexec -np 4 ./a.out

info
```
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);
```

Example

```
int main(int argc, char **argv) {
  MPI_Init(&argc, &argv);
  int num;
  MPI_Comm_rank(MPI_COMM_WORLD, &num);
  if(num == 0) { //master
    MPI_status status;
    char txt[100];
    MPI_Recv(txt, 100, MPI_CHAR, 1, 42,
             MPI_COMM_WORLD, &status);
    std::cout << txt << '\n';
  } else { //worker
    std::string text = "Bla";
    MPI_Send(const_cast<char*>(text.c_str()),
             text.size()+1, MPI_CHAR, 0, 42,
             MPI_COMM_WORLD);
  }
  MPI_Finalize();
}
```

Open MP

compile g++ -fopenmp -std=c++11

import #include <omp.h>

info
```
omp_get_thread_num()   → id
omp_get_num_threads()  → #threads
```

#PCSE I HS16

F. Vogel

Parallel Scaling

Amdahls Law: $S(n) = \frac{1}{(1-p) + \frac{p}{n}}$

Gustafsons Law: $S(n) = 1 - p + n \cdot p$ (weak scaling)

Strong scaling: $S(n) = \frac{T(1)}{T(n)}$ $E(n) = \frac{S(n)}{n}$

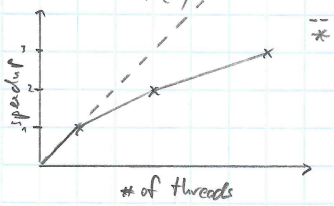

Weak scaling: $E(n) = \frac{T(1)}{T(n)}$ fixed problem size per thread

Note: Problem size for N^2 solver is N^2 not N

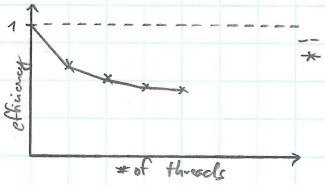

Roofline Model

Operational Intensity [Flop/Byte] r

→ $r = \frac{\#\text{ of Flops}}{\#\text{ of Mem accesses} \cdot \#\text{ of Bytes}}$

→ Example: $RHS_{i,j} = C_n \cdot (q^n_{i-1,j} + q^n_{i-1,j} + q^n_{i,j+1} + q^n_{i,j-1} - 4q^n_{i,j})$

→ 6 Flops (4 ADD + 2 MUL)

→ Mem Access:

↦ No cache: 5 read + 1 write = 6

↦ Infinite cache: 1 read + 1 write = 6

→ Op. Int.

↦ No cache: $\frac{6\text{Flops}}{6 \cdot 4\text{B}} = 0.25$ Flop/B

↦ Infinite cache: $\frac{6\text{ Flops}}{2 \cdot 4\text{B}} = 0.75$ Flop/B

Nominal performance

→ f_{peak} = processor clock/sec · vector size · instructions/clock · # cores

→ b_{peak} = memory clock/sec · channel size · #channels / 8[bit/Byte]

Plot

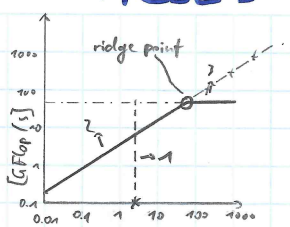

1: Locality

2: Communication

3: Computation

limit by comp. ---- $f_{peak} = f(r_k)$

limit by DRAM -·-·- $r_k \cdot b_{peak} = f(r_k)$

Bugs

→ Threads : Pass thread id by value, not by reference

→ SIMD : Race condition by interleaving of loads

→ SIMD : Undefined behavior by loading not allocated memory

→ OMP : implicit barrier @ end of for, all threads have to call it. Fix: schedule(dynamic) nowait

→ MPI : Don't update asynchronous send buffer

→ Threads : Not passing all needed variables

→ Threads : Race condition by not locking result double

→ Threads : Result reduction before join is called on each thread

→ Threads : Master thread: not releasing lock before calling join can result in deadlock

→ MPI : Ordering needed to prevent deadlock (replacing with asynchronous works too)